

Banin Abrar

✉ babrar@uwaterloo.ca — 📞 519-729-6017 — [in/baninabrar](#) — 🌐 babrar — 🌐 babrar.xyz

SKILLS

- **Languages & Frameworks:** Python, Rust, C/C++, Golang, TypeScript, Java, SQL; React, Node, Django, GraphQL
- **Tools:** Kubernetes, Terraform, AWS, GCP, gRPC, Kafka, Spark, Prestodb, Elasticsearch, dbt, PostgreSQL, BigQuery

EDUCATION

University of Waterloo

Class of 2023

Bachelor of Applied Science in Computer Engineering

- **Relevant Courses:** Programming for Performance (*Rust*), Distributed Systems, Networks & Security, Operating Systems, Compilers, Database Systems, Software Design and Architectures, Data Structures and Algorithms

WORK EXPERIENCE

Sony PlayStation

San Diego, CA

Software Engineering Intern

Sep 2022 – Dec 2022

- Resolved critical performance bottleneck in the replication of PlayStation Plus data for **50 million** users
- Developed Ceph-powered K8s operator to rapidly provision test environments with prod data to fix the bottleneck
- Enhanced workflow efficiency for 30+ engineers by creating a React frontend for database configuration tools

Arctic Wolf

Waterloo, ON

Software Developer Co-op, Infrastructure

Jan 2022 – Apr 2022

- Reduced AWS performance alerts for ElasticSearch by **50%** and prevented high CPU loads during peak hours
- Created and deployed a Go-based watchdog to optimize index-sharding and attain aforementioned goal
- Integrated Datadog metrics into Grafana to enhance microservice monitoring and enable faster incident responses

Sheertex

Montreal, QC

Software Engineering Co-op, Business Intelligence

Jan 2021 – Aug 2021

- Implemented a modern ELT process using Snowflake, dbt and Fivetran, delivering significantly more accurate KPIs to stakeholders and saving data analysts over **10 hours/week** on BI workflows
- Streamlined inventory data access by utilizing Google Cloud Functions to load data into the main data warehouse

Loblaw Digital

Toronto, ON

Site Reliability Engineer

May 2020 – Aug 2020

- Successfully migrated logging infrastructure from Splunk to Google's Stackdriver platform using Terraform and BigQuery SQL scripts, netting over **\$50K** in annual savings, while incurring minimal performance penalties
- Developed an internal library for Google PubSub to seamlessly produce and consume data in a specialized format

FLIR Systems

Waterloo, ON

Tools Developer Co-op

May 2019 – Aug 2019

- Enhanced performance of S3 queries by implementing a highly performant data stack using Spark, Presto, and Hive

JANA Solutions

Aurora, ON

Junior Software Developer

Sept 2018 – Dec 2018

- Delivered a comprehensive full-stack application written in Django and React for statistical modelling to **10+** clients

OTHER EXPERIENCE & PROJECTS

Waterloop Student Design Team (*TypeScript, React, MongoDB*)

University of Waterloo

- Revamped a static web page into a dynamic MERN stack website with enhanced features, optimizing user experience

1nil (*Python, gRPC, PostgreSQL*)

[fydp-1nil/backend](#) 🌐

- Open-source video streaming and data analysis platform for minor and amateur sports leagues

FastOLAP (*PrestoDB, Hive, Redash*)

[babrar/fastolap](#) 🌐

- An open-source fork of AWS Athena for performing fast analytic SQL queries on large datasets hosted on Amazon S3